


Design and Optimization of Mobile Aggregation Architectures

Ioannis Papapanagiotou, PhD
(Yanni)

Computer & Information Technology
Purdue University



Who am I?

- Dual Major PhD in North Carolina State University
 - Computer Engineering
 - Operations Research
- IBM (3 years)
 - Emerging Technology Institute:
 - In house incubator team reporting directly to the CTO and VP
 - TJ Watson Research Center:
 - IT & Wireless convergence group
- NCSU (3 years RA & TA)
 - Taught Engineering, Networking and Simulation courses
 - Research funded by IBM, Cisco Systems, Time Warner Cable
- Appointments with
 - University of Patras
 - Technical Chamber of Greece
 - VTT Technical Institute in Finland
- Fellowships:
 - IBM PhD Fellow
 - Academy of Athens PhD Fellow

What I do

- Current work
 - Server side NoSQL/JSON solution for end-stores
 - Working with front-end caches and databases
- Research Interests:
 - Data analytics in wireless networks
 - Byte level deduplication
 - Performance analysis & simulations
 - Elasticity in Cloud Computing
- Experience with E-Learning technologies:
 - Synchronous (e.g. virtual meeting rooms) and Asynchronous (e.g. video recordings) types of assisting.
 - Emulation of computer networking labs through virtualization.
- Publication Record
 - 5 Peer-Reviewed journals
 - 12 Conference proceedings
 - 8 Submitted patents
 - 3 Book chapters
 - 121 Citations
- More info: <http://people.engr.ncsu.edu/ipapapa>


Motivation

Smartphones



+ =

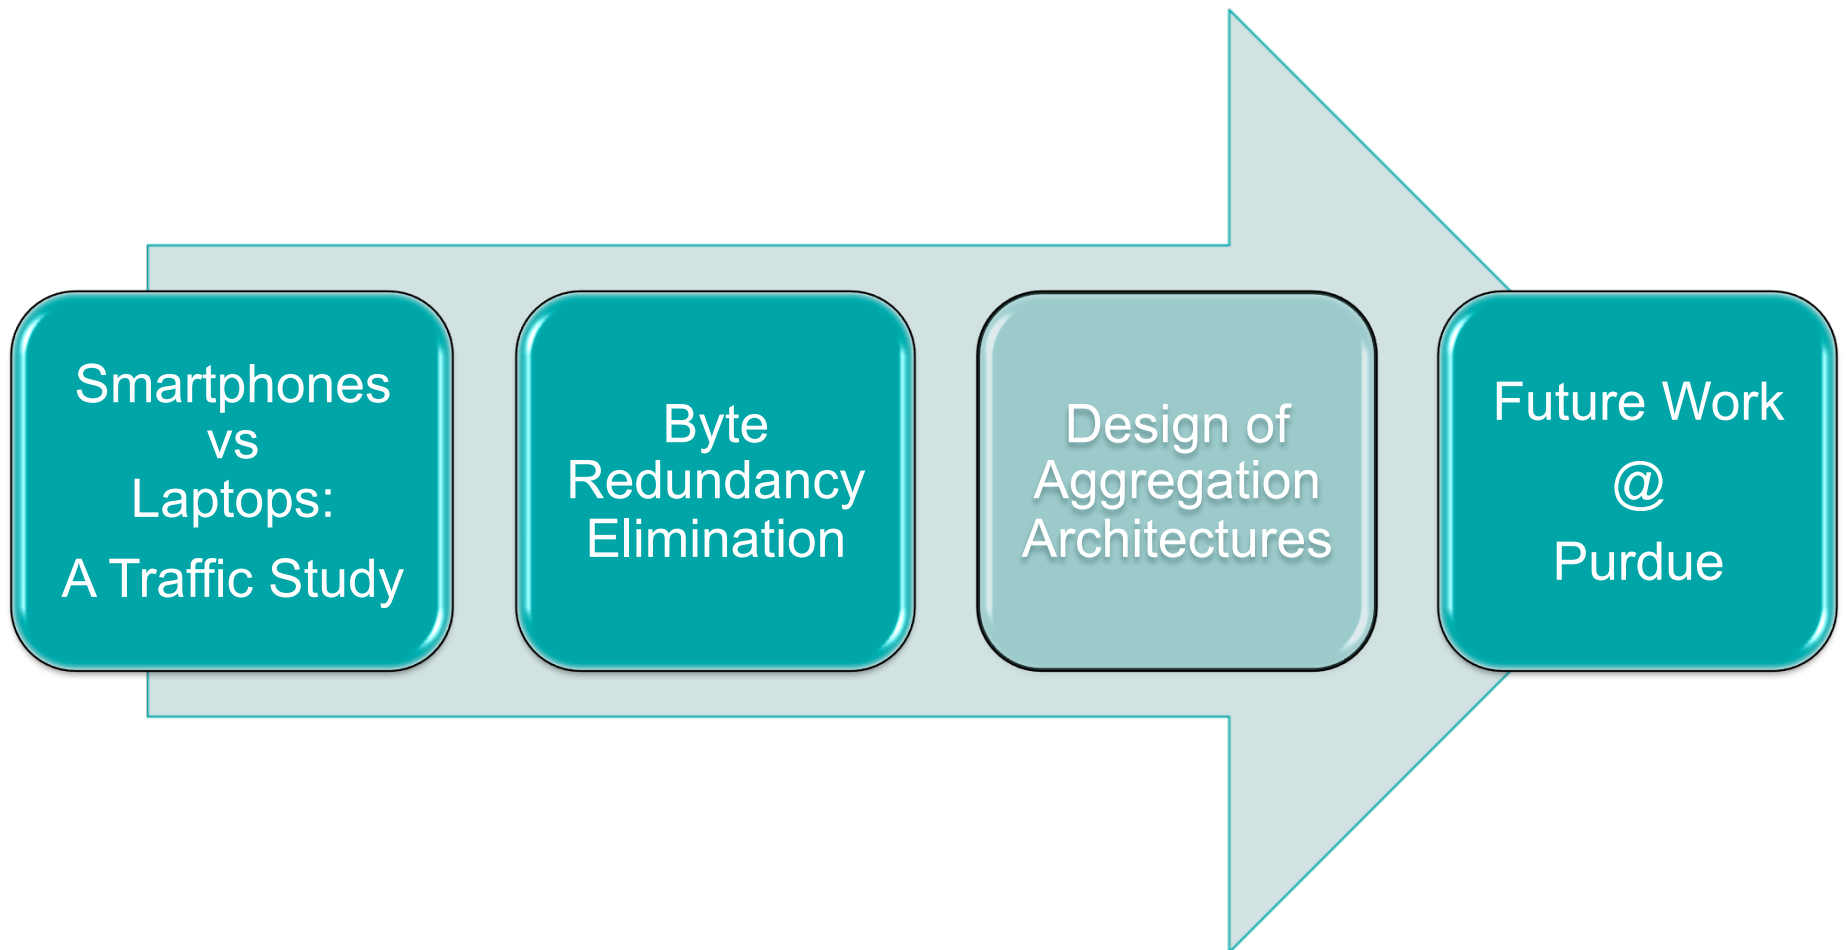
Need For Speed



Wireless Service Providers have to transfer this traffic to the Internet!

- **Operating System fingerprinting** algorithm based on data mining principles
 - Compared Smartphones vs Laptops.
- Analyzed **IP lease traffic behavior**
 - Proposed **novel IP lease policy** that improves address space utilization six-fold and without increasing the overhead.
- Analyzed **Web traffic behavior** and object caching
 - Demonstrated that a 10MB browser cache can save on average ~5%.
 - Showed the importance of **storing policies in proxy caches**.
- Designed a “**Hybrid**” **Byte Redundancy Elimination**
 - Provides more savings, with less memory space and is x3 faster.
 - Added QoS in redundancy elimination.


Presentation Outline



Section 1

Smartphones vs Laptops: A traffic study

Trace Collection



Trace Type	IBM Research	NCSU Centennial
Dates (2012)	Feb 29 – Mar 25	Jan 15 – Feb 15
Client MAC address	2980	8726
Available IP address	4096 (8 * /23)	2048 (/21)
DHCP Leases	12h	15 min
TCP/UDP Bytes	2.5TB	4.9TB
Software Used	Hacked Bro IDS 2.0	

Device Identification

- We found correlation between the fields of the DHCP Request header and the OS
 1. *Host-Name*
 2. *Vendor-Name*
 3. *Parameter-Request-List*
 4. Organization Unique Identifier (OUI)
 5. *Options* parameter sequence
- Developed a **data mining algorithm** that quantifies the correlation

```
Frame 75: 342 bytes on wire (2736 bits), 342 bytes captured (2736 bits) on interface 0
Ethernet II, Src: 42:09:02:fa:03:00 (42:09:02:fa:03:00), Dst: [redacted]
Internet Protocol Version 4, Src: [redacted], Dst: [redacted]
User Datagram Protocol, Src Port: bootps (67), Dst Port: bootps (67)
Bootstrap Protocol
  Message type: Boot Request (1)
  Hardware type: Ethernet
  Hardware address length: 6
  Hops: 1
  Transaction ID: 0xa57f8cd9
  Seconds elapsed: 0
  Bootp flags: 0x0000 (Unicast)
  Client IP address: 0.0.0.0 (0.0.0.0)
  Your (client) IP address: 0.0.0.0 (0.0.0.0)
  Next server IP address: 0.0.0.0 (0.0.0.0)
  Relay agent IP address: 9.2.240.4 (9.2.240.4)
  Client MAC address: Apple_92:f4:73 (00:26:b0:92:f4:73)
  Client hardware address padding: 00000000000000000000
  Server host name not given
  Boot file name not given
  Magic cookie: DHCP
  Option: (t=53,l=1) DHCP Message Type = DHCP Request
  Option: (t=55,l=6) Parameter Request List
    Option: (55) Parameter Request List
    Length: 6
    Value: 0103060f77fc
    1 - Subnet Mask
    3 - Router
    6 - Domain Name Server
    15 - Domain Name
    119 - Domain Search [TODO:RFC3397]
    252 - Private/Proxy autodiscovery
  Option: (t=57,l=2) Maximum DHCP Message Size = 1500
  Option: (t=61,l=7) Client identifier
  Option: (t=50,l=4) Requested IP Address = [redacted]
  Option: (t=51,l=4) IP Address Lease Time = 90 days
  Option: (t=12,l=6) Host Name = "iPhone"
  End option
  Padding
```

First 3 Bytes of MAC Address

"Parameter Request List" fields sequence

"Options" fields sequence

Host Name

Classification Outcome


		IBM Research		NCSU Centennial	
Device	OS	#	%	#	%
Laptop	All	2176	73.02	3970	45.50
	Windows	1787	59.92	2819	32.31
	Mac OS X	385	12.92	1131	12.96
	Linux	4	0.13	20	0.23
Smartphone	All	735	23.66	4489	51.44
	iPhone/iPad/iPod	577	19.36	3069	35.17
	Android	126	4.24	1336	15.29
	BlackBerry	31	1.04	84	0.96
	Windows Mobile	1	0.03	2	0.02
Other	All	69	2.32	267	3.06
	Cisco VoIP	9	0.32	-	-
	Unidentified	60	2.01	267	3.06
All		2980	100	8726	100

Wrote a trace-driven simulator to evaluate different lease policies:

- a) Static Policies:** Fixed lease of 15 min, 30 min or 12 hours (most common case).
- b) Exponential Adaptation:** Allocates a short lease to client once it arrives, and doubles the lease time every time the client renews the lease [GaTech IMC 2007]
- c) Differential Lease:** Allocates different lease values based on device:


Trace	iOS	Android	RIM	Windows	MAC
Corporate	1000	2000	2000	4000	4000
Educational	500	1000	1000	2000	2000

Simulation Results




- Differential lease policy performs well:
 - Achieves low address space utilization (comparable to short lease time)
 - Reasonable number of broadcasts (compared to short lease times)
 - Reasonable server load

Web Traffic: Response Code



- An HTTP Response code indicates the “status” of the object in the server.
- **304 Not Modified:** validate the client’s cached copy on the browser cache.
 - Windows and Mac OS X generate more than smartphones.
- **206 Partial Content:** used for large objects that must be chunked.
 - iOS generate way more than any other OS.

Web Traffic: Proxy Cache Advantages & Metrics





- **Web proxy caches** are used to improve the performance of the Web.
 - Reduce network bandwidth
 - Reduce web server load
 - Improve response time
- **Request Hit Rate:** The percentage of all requests that can be satisfied by searching the cache for a copy of the requested object.
- **Byte Hit Rate:** The percentage of all data that is directly transferred from the cache, rather than the origin server.

Web Traffic: Proxy Cache Size & Freshness



We evaluate the effectiveness based on two configurations for objects that are not fresh in cache:

1. *Never Revalidated*: The cached object is the same as the one in the server.
2. *Always Revalidated*: The cached object is different than the one in the server.



- Diminishing returns ~100GB, hence replacement policies are not relevant any more.
- Savings range from 10-25%, and Request hit rates 10-15%.

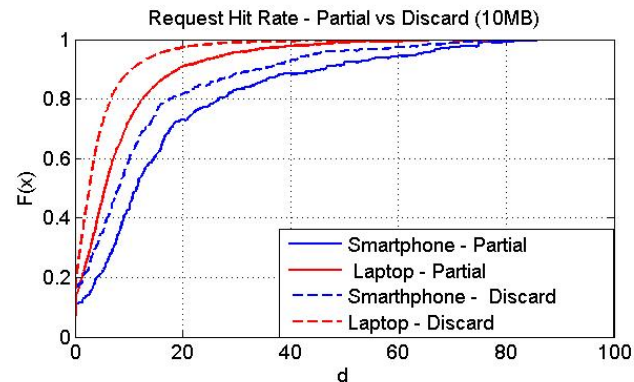


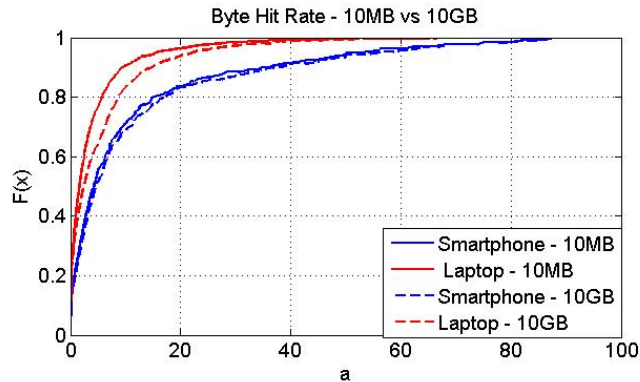
Web Traffic: Proxy Cache Storing Policies



- We evaluate the case where the cache is configured to download X% more data than the data requested by the client.
- **Full download** (cache pre-fetches the whole object) policy results in abusive BW demands in smartphones.
- **Partial storing** (cache stores only what user requested) is the optimum policy.
- A partially downloaded object is cached in its entirety if more than X% of its size has been downloaded by the client.
- 0% corresponds to the **Full download** policy.
- **Discarding** (any partially downloaded object is not cached) is more beneficial than other conditional download policies.

Web Traffic: Browser Cache Efficiency

We replay the IBM trace for each device through an additional simulated cache dedicated solely to that device:



- Current laptop browser caches are more efficient than the smartphone ones.
- A small browser cache (10MB) is sufficient to capture most of the savings in smartphones.
- A browser cache should be able to handle partially downloaded objects.

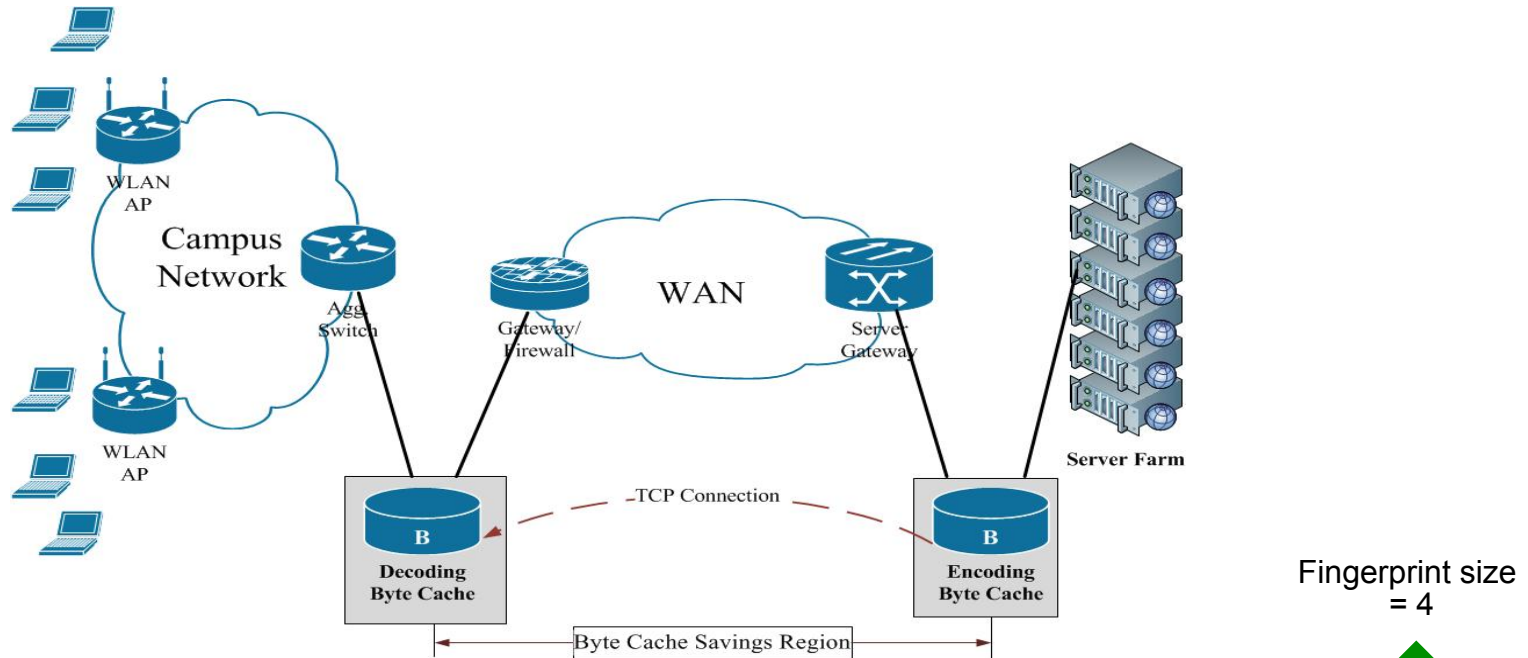
Section 2

Byte Redundancy Elimination in Wireless Networks

Proxy Caches:

- Limited savings: Partially modified objects are treated as different ones.
- Almost 60% of the Web data are uncacheable (based on RFC2616).
- From the cacheable objects, almost 70% may be referenced only once.
- Object revalidation may result in fetching the object from the server.

Byte Cache: Introduction



A Byte Cache is a system of devices that can identify redundancy at the byte level.

- Requires two middleware boxes (encoder & decoder).
- Uses Rabin fingerprinting to generate the hashes.
- Hashes are then used as indices to the content, or to break the content in chunks.

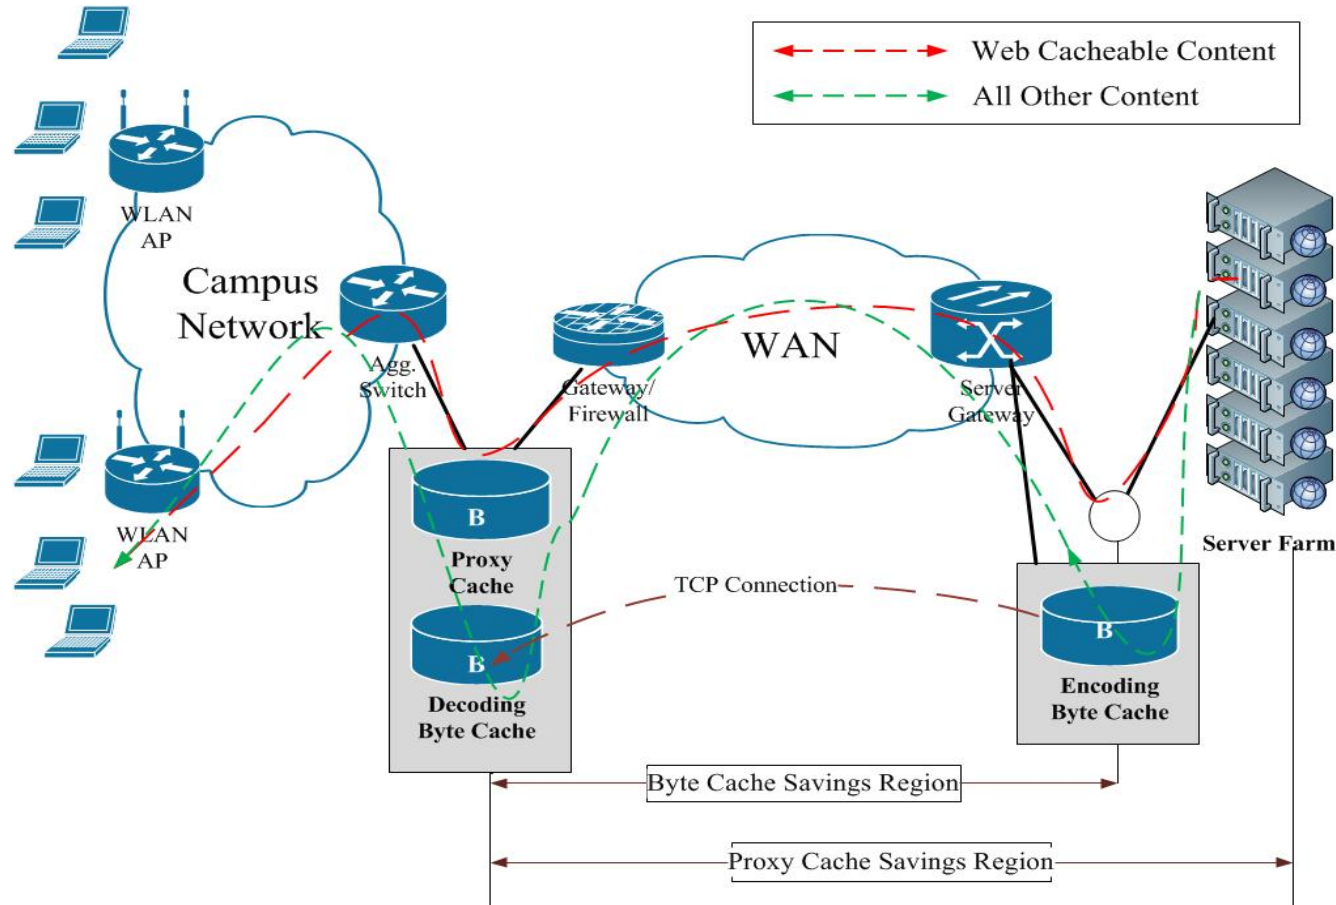
HDD	RAM
<u>Content</u>	<u>Hash</u>
ABCDEFEGHIJ	H(ABCD) H(BCDE) H(CDEF) H(DEFG) H(EFGH) H(FGHI) H(GHIJ)

Fingerprint size
= 4

Packet 1
srcIP: 15.2.1.1
dstIP: 150.150.1.1
Payload: ABCDEFEGHIJ

- Fingerprint generation is CPU intensive.
 - Can be a bottleneck in high-bandwidth WAN links.
- Requires a lot of memory space.
 - A hash index is required per object in proxy caches vs. per 32B chunks in DRE systems
- Fingerprint generation and chunk storage may be useless for some protocols.
 - E.g. Encrypted traffic.


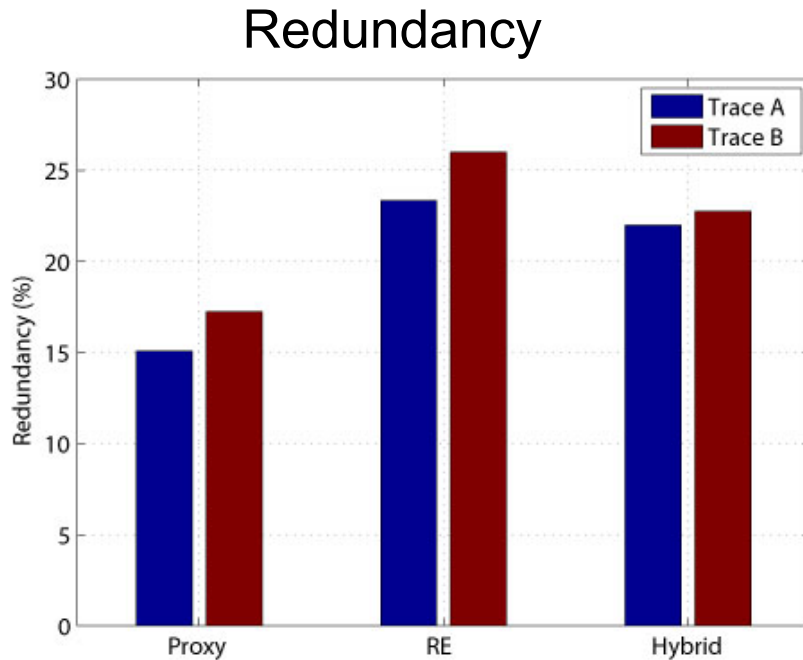
Hybrid Byte Cache: Static Scheduler



- Scheduler:
 - Web cacheable content is forwarded to a **Proxy Cache** Module
 - Web uncacheable content is forwarded to an **Encoding Byte Cache** Module
 - Other data are sent unprocessed

Hybrid Byte Cache: Redundancy & Savings

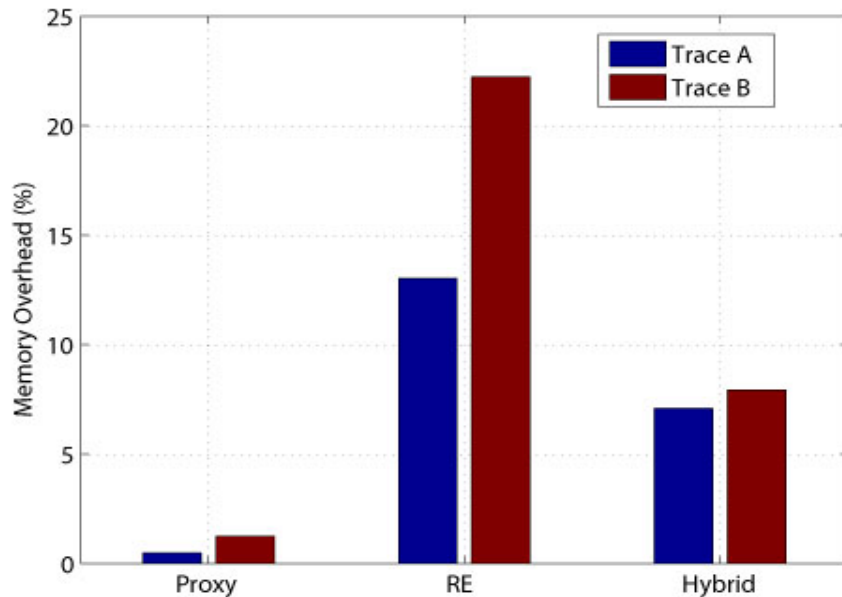
- Full packet traces captured in IBM (Trace A: 19GB, Trace B: 64GB)
 - Both traces consist of 80% TCP and 75% of Web traffic.



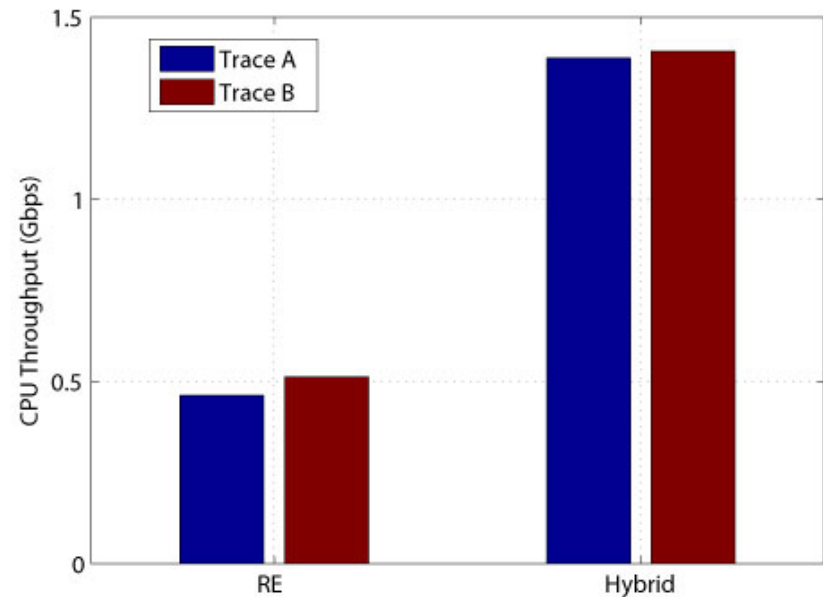
- Redundancy in DRE is better than in Proxy or Hybrid DRE
 - Note that redundancy includes the hash overhead.
 - Providers are mostly interested for Savings.
- Hybrid DRE provides more savings than Proxy or DRE.

Hybrid Byte Cache: System Performance

Memory requirement



CPU Throughput



- “Memory Overhead” indicates the amount of RAM needed to store the hashes that correspond to the content.
- Hybrid DRE requires more than half the memory compared to DRE.
- Hybrid DRE increases throughput almost 3 times compared to DRE.
 - Therefore, Hybrid DRE can be installed in faster links.

Conclusion

Summary

- Differentiated DHCP Lease policy should be implemented in wireless networks.
- A Smartphone vendor should consider adding an extra **10MB** of Browser Cache:
 - ~5% of savings.
- A **100GB**, properly configured proxy cache can provide:
 - 25% of savings for smartphones.
 - 15% of savings in a wireless network.
- A Byte Cache system may provide more savings than a proxy.
- The proposed Hybrid DRE leverages the advantages of both the Proxy Cache and the Byte Cache:
 - Better savings
 - Less CPU cycles
 - Less memory

Future Work at Purdue University

■ Collaborations


- IBM Research, Cisco Systems and Time Warner
- North Carolina State University and University of Michigan
 - NSF proposals
- European Universities

■ Research Areas

- Data analytics in mobile networks
- Byte level deduplication
- Performance analysis & simulations
- Elasticity in Cloud Computing
- Applications to the Smart Grid

■ Extend Intellectual Property


Future Work #1: Dynamic Hybrid DRE



Dynamic Scheduler:

- Based on machine learning, predicts the best caching module for each traffic flow.
- Optimized to take into account the savings and system performance.
- Decreases CPU usage, benefits from parallel processing, increases savings and system failover.

Future Work #2: Quality of Service in Data Deduplication



- Users are differentiated in class categories:
 - Each class processes the flows with a different set of parameters.
- We introduce a controlling mechanism that dynamically allocates memory and space:
 - Per-class controllers monitor the performance and send messages to the central controller such that the SLAs are met.

Thank you!